

# SQL Stored Procedures

Winter 2006-2007

Lecture 10

# SQL Functions

- SQL queries can use sophisticated math operations and functions
  - Can compute simple functions, aggregates
  - Can compute and filter results
- Applications often require specialized computations
  - Would like to use these in SQL queries, too
- SQL provides a mechanism for defining functions
  - Called User-Defined Functions (UDFs)

# SQL Functions (2)

- Can be defined in a procedural SQL language, or in an external language
  - SQL:1999, SQL:2003 both specify a language for declaring functions and procedures
  - Different vendors provide their own languages
    - Oracle: PL/SQL
    - Microsoft: TransactSQL
    - PostgreSQL: PL/pgSQL
    - MySQL: new stored procedure support strives to follow specifications (and mostly does)
    - Some also support external languages: Java, C, etc.
  - Lots of variation in features and support (as usual)

# Example SQL Function

- A SQL function to count how many bank accounts a particular customer has:

```
CREATE FUNCTION
    account_count(customer_name VARCHAR(20))
RETURNS INTEGER
BEGIN
    DECLARE a_count INTEGER;
    SELECT COUNT(*) INTO a_count
        FROM depositor
        WHERE depositor.customer_name = customer_name;
    RETURN a_count;
END
```

- Function can take arguments and return values
- Use SQL statements and other operations in body

# Example SQL Function (2)

- Can use our function for individual accounts:

```
SELECT account_count( 'Johnson' );
```

- Can include in computed results:

```
SELECT customer_name,  
       account_count(customer_name) AS accounts  
FROM customer;
```

- Can include in **WHERE** clause:

```
SELECT customer_name  
FROM customer  
WHERE account_count(customer_name) > 1;
```

# Arguments and Return-Values

- Functions can take any number of arguments
- Functions *must* return a value
  - Specify type of value in **RETURNS** clause
- From our example:

```
CREATE FUNCTION
```

```
    account_count(customer_name VARCHAR(20))
```

```
    RETURNS INTEGER
```

- One argument named **customer\_name**, type is **VARCHAR(20)**
- Returns some **INTEGER** value

# Table Functions

- SQL:2003 spec. includes table functions
  - Return a whole table as their result
  - Can be used in **FROM** clause
- A generalization of views
  - Can be considered to be parameterized views
  - Call function with specific arguments
  - Result is a relation based on those arguments
- Although SQL:2003 not broadly supported yet, most DBMSes provide a feature like this

# Function Bodies and Variables

- Blocks of SQL commands are enclosed with **BEGIN** and **END**
  - Defines a compound statement
  - Can have nested **BEGIN ... END** blocks
- Variables are specified with **DECLARE** command
  - Must appear at start of block
  - Initial value is **NULL**
  - Can initialize to some other value with **DEFAULT** syntax
  - Scope of a variable is within its block
  - Variables in inner blocks can shadow variables in outer blocks



# Example Blocks and Variables

- Our `account_count` function's body:

```
BEGIN
    DECLARE a_count INTEGER;
    SELECT COUNT(*) INTO a_count
        FROM depositor
        WHERE depositor.customer_name = customer_name;
    RETURN a_count;
END
```

- A simple integer variable with initial value:

```
BEGIN
    DECLARE result INTEGER DEFAULT 0;
    ...
END
```

# Assigning To Variables

- Can use **SELECT ... INTO** syntax
  - For assigning the result of a query into a variable

```
SELECT COUNT(*) INTO a_count
FROM depositor
WHERE depositor.customer_name=customer_name;
```
  - Query must produce a single row
  - Note:** **SELECT INTO** often has multiple meanings!  
This form is specific to stored routines
- Can also use **SET** syntax
  - For assigning result of a math expression to a variable

```
SET result = n * (n + 1) / 2;
```

# Assigning Multiple Variables

- Can assign to multiple variables using **SELECT INTO** syntax
- Example: Want both the number of accounts and the total balance

```
DECLARE a_count INTEGER;  
DECLARE total_balance NUMERIC(12,2);  
SELECT COUNT(*), SUM(balance)  
    INTO a_count, total_balance  
    FROM depositor NATURAL JOIN account  
    WHERE depositor.customer_name=customer_name;
```

# Another Example

- Simple function to compute sum of 1..N

```
CREATE FUNCTION sum_n(n INTEGER) RETURNS INTEGER
BEGIN
    DECLARE result INTEGER DEFAULT 0;
    SET result = n * (n + 1) / 2;
    RETURN result;
END
```

- Lots of extra work in that. To simplify:

```
CREATE FUNCTION sum_n(n INTEGER) RETURNS INTEGER
BEGIN
    RETURN n * (n + 1) / 2;
END
```

# Dropping Functions

- Can't simply overwrite functions in the database
  - Same as tables
- First, drop old version of function:

```
DROP FUNCTION sum_n;
```

- Then create new version of function:

```
CREATE FUNCTION sum_n(n INTEGER)
RETURNS INTEGER
BEGIN
    RETURN n * (n + 1) / 2;
END
```

# SQL Procedures

- Functions have specific limitations
  - Must return a value
  - All arguments are input-only
  - Typically cannot affect transaction status (i.e. commit, rollback, etc.)
  - Usually not allowed to modify relations, except in particular circumstances
- Procedures are more general constructs without these limitations
  - Also generally cannot be used in same places as functions

# Example Procedure

- Write a procedure that returns both the number of accounts a customer has, and their total balance

- Results are passed back using out-parameters

```
CREATE PROCEDURE account_summary(  
    IN customer_name VARCHAR(20),  
    OUT a_count INTEGER,  
    OUT total_balance NUMERIC(12,2))  
BEGIN  
    SELECT COUNT(*), SUM(balance)  
        INTO a_count, total_balance  
        FROM depositor NATURAL JOIN account  
        WHERE depositor.customer_name = customer_name;  
END
```

- Default parameter type is IN

# Calling a Procedure

- Use the **CALL** statement to invoke a procedure

```
CALL account_summary(...);
```

- To use this procedure, must also have variables to receive the values

- MySQL SQL syntax:

```
CALL account_summary('Johnson',  
                      @j_count, @j_total);
```

```
SELECT @j_count, @j_total;
```

- **@var** declares a temporary session variable

@j_cnt	@j_tot
2	1400.00



# Conditional Operations

- SQL provides an if-then-else construct

```
IF cond1 THEN command1  
ELSE IF cond2 THEN command2  
ELSE command3  
END IF
```

- Branches can specify compound statements instead of single statements
  - Enclosed with **BEGIN** and **END**
- Can leave out **ELSE IF** and/or **ELSE** clauses

# Looping Constructs

- SQL also provides looping constructs
- **WHILE** loop:

```
DECLARE n INTEGER DEFAULT 0;  
WHILE n < 10 DO  
    SET n = n + 1;  
END WHILE;
```

- **REPEAT** loop:

```
REPEAT  
    SET n = n - 1;  
UNTIL n = 0  
END REPEAT;
```

# Iteration Over Query Results

- Sometimes need to issue a query, then iterate over each row in result
  - Perform more sophisticated operations than simple SQL query can perform
- Example:
  - Write a function that handles customer withdrawals in a special way:
    - Money is withdrawn from first customer account.
    - If first account runs out of money, move to customer's second account and withdraw more.
    - Repeat until the total withdrawal request is satisfied, or the customer's accounts run out of money.
  - Need to operate on each row of customer's accounts

# Cursors

- Need to issue a query to fetch specific results
- Need to iterate through each row of results
  - Operate on each row's values individually
- A cursor is an iterator over rows in a result set
  - Cursor refers to one row in query results
  - Can access row's values through cursor
  - Can move cursor forward through results
- Cursors can provide different features
  - Read-only vs. read-write
  - Forward-only vs. bidirectional
  - Static vs. dynamic (when concurrent changes occur)

# Cursor Notes

- Cursors can be expensive
- Can operation use a normal SQL query instead?
  - (Usually, yes...)
  - Cursors let you do what databases do, slowly
- Cursors might hold system resources until they are finished
  - e.g. database might store query results in a temp table, to present read-only, static view of query result
- Syntax varies widely across DBMSes
- Most external DB connectivity APIs provide cursor capabilities

# Stored Procedures and Cursors

- Can use cursors inside stored procedures
- Book syntax:

```
DECLARE n INTEGER DEFAULT 0;  
FOR r AS SELECT balance FROM account  
          WHERE branch_name='Perryridge'  
DO  
    SET n = n + r.balance;  
END FOR
```

- Iterates over account balances from Perryridge branch, summing balances
- (Could write this with a simple SQL query, too)

# MySQL Cursor Syntax

- Must *explicitly* declare cursor variable

```
DECLARE cur CURSOR FOR  
SELECT ... ;
```

- Open cursor to use query results:

```
OPEN cur;
```

- Fetch values from cursor into variables

```
FETCH cur INTO var1, var2, ... ;
```

- Next row is fetched, values stored into specified variables
- Must specify same number of variables as columns in result
- A specific error condition is flagged to indicate end of results

- Close cursor at end of operation

```
CLOSE cur;
```

- Also happens automatically at end of block

# Handling Errors

- Many situations where errors can occur in stored procedures
  - Called conditions
  - Includes errors, warnings, other signals
  - Can also include user-defined conditions
- Handlers can be defined for conditions
- When a condition is signaled, its handler is invoked
  - Handler can specify whether to continue running the procedure, or whether to exit procedure instead



# Conditions

- Predefined conditions:
  - **NOT FOUND**
    - Query fetched no results, or command processed no results
  - **SQLWARNING**
    - Non-fatal SQL problem occurred
  - **SQLException**
    - Serious SQL error occurred

# Conditions (2)

- Can also define application-specific conditions
  - Examples:
    - “Account overdraft!”
    - “Inventory of item hit zero.”
- Syntax for declaring conditions:

```
DECLARE acct_overdraft CONDITION
DECLARE zero_inventory CONDITION
```
- Not every DBMS supports generic conditions
  - e.g. MySQL supports assigning names to existing SQL error codes, but not general conditions

# Handlers

- Can declare handlers for specific conditions
- Handler specifies statements to execute
- Handler also specifies what happens next:
  - Continue running procedure where left off
  - Exit stored procedure completely
- Syntax:
  - A continue-handler:  
`DECLARE CONTINUE HANDLER FOR cond statement`
  - An exit-handler:  
`DECLARE EXIT HANDLER FOR cond statement`
  - Can also specify a statement-block

# Handlers (2)

- Handlers can do very simple things
  - e.g. set a flag to indicate some situation
- Can also do very complicated things
  - e.g. insert rows into other tables to log failure situations
  - e.g. properly handle an overdrawn account

# Total Account Balance – MySQL

- Declared as a function – returns a value

```
CREATE FUNCTION acct_total(cust_name VARCHAR(20))
RETURNS NUMERIC(12,2)
BEGIN
    -- Variables to accumulate into
    DECLARE bal NUMERIC(12,2);
    DECLARE total NUMERIC(12,2) DEFAULT 0;

    -- Cursor, and flag for when fetching is done
    DECLARE done INT DEFAULT 0;
    DECLARE cur CURSOR FOR SELECT balance
        FROM account NATURAL JOIN depositor
        WHERE depositor.customer_name = cust_name;
```

# Total Account Balance (2)

```
-- When fetch is complete, handler sets flag
DECLARE CONTINUE HANDLER FOR SQLSTATE '02000'
    SET done = 1;

OPEN cur;
REPEAT
    FETCH cur INTO bal;
    IF NOT done THEN
        SET total = total + bal;
    END IF;
UNTIL done END REPEAT;
CLOSE cur;
RETURN total;
END
```

# Using Our Stored Procedure

- Can compute total balances now:

```
SELECT customer_name,  
       acct_total(customer_name) AS total  
FROM customer;
```

– Result:

customer_name	total
Adams	0.00
Brooks	0.00
Curry	0.00
Glenn	0.00
Green	0.00
Hayes	900.00
Jackson	0.00
Johnson	1400.00
Jones	750.00
Lindsay	700.00
Majeris	850.00
McBride	0.00
Smith	1325.00
Turner	350.00
Williams	0.00

# Stored Procedure Benefits

- Very effective for manipulating large datasets within the database
  - Don't incur communications overhead
  - Database is designed for such tasks, but applications and languages typically aren't
- Often used to provide secure interface to data
  - Banks will lock down data, and only expose operations through stored procedures
- Can encapsulate business logic in procedures
  - Forbid invalid states by requiring all operations go through stored procedures



# Stored Procedure Drawbacks

- Increases load on database system
  - Can reduce performance for *all* operations being performed by DBMS
  - Need to make sure the operation *really* requires a stored procedure...
- Hard to migrate to a different DBMS
  - Different vendors' procedural languages have *many* distinct features and limitations